

BC_DT_SVG

August 15, 2020

```
[61]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn import metrics
import matplotlib.pyplot as plt
from sklearn import tree
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
```

```
[62]: #Import data file: Breast Cancer
BC_df = pd.read_csv("C:/Venu/UCI DataSets/breast-cancer.data",delimiter=',',
→header=None,names=['Class', 'Age', 'Menopause', 'Tumor-size', 'Inv-nodes',
→'Node-caps', 'Deg-malig', 'Breast', 'Breast-quad', 'IR-Radiat'])
```

```
[63]: #Checking for Missing Values/Null Values
BC_df.isna().values.any()
BC_df.isna().sum()
BC_df.isnull()
BC_df.isnull().sum()
```

```
[63]: Class          0
Age              0
Menopause        0
Tumor-size       0
Inv-nodes        0
Node-caps        8
Deg-malig        0
Breast           0
Breast-quad      1
IR-Radiat        0
dtype: int64
```

```
[64]: #Removing the Missing rows
BC1_df = BC_df.dropna()
BC1_df.isna().sum()
```

```
[64]: Class      0
      Age        0
      Menopause  0
      Tumor-size 0
      Inv-nodes  0
      Node-caps  0
      Deg-malig  0
      Breast     0
      Breast-quad 0
      IR-Radiat  0
      dtype: int64
```

```
[65]: #Response variable counts
      BC1_df.Class.value_counts()
```

```
[65]: no-recurrence-events    196
      recurrence-events       81
      Name: Class, dtype: int64
```

```
[66]: #Converting the categorical features as dummy variables
      encoded_BC1_df = pd.get_dummies(BC1_df,drop_first=True)
      X_features = list(encoded_BC1_df.columns)
      X_features
      X_features.remove('Class_recurrence-events')
      X_features
      encoded_BC1_df[X_features].columns
```

```
[66]: Index(['Deg-malig', 'Age_30-39', 'Age_40-49', 'Age_50-59', 'Age_60-69',
            'Age_70-79', 'Menopause_lt40', 'Menopause_premeno', 'Tumor-size_10-14',
            'Tumor-size_15-19', 'Tumor-size_20-24', 'Tumor-size_25-29',
            'Tumor-size_30-34', 'Tumor-size_35-39', 'Tumor-size_40-44',
            'Tumor-size_45-49', 'Tumor-size_5-9', 'Tumor-size_50-54',
            'Inv-nodes_12-14', 'Inv-nodes_15-17', 'Inv-nodes_24-26',
            'Inv-nodes_3-5', 'Inv-nodes_6-8', 'Inv-nodes_9-11', 'Node-caps_yes',
            'Breast_right', 'Breast-quad_left_low', 'Breast-quad_left_up',
            'Breast-quad_right_low', 'Breast-quad_right_up', 'IR-Radiat_yes'],
            dtype='object')
```

```
[67]: #Final X/Predictors set
      X = encoded_BC1_df[X_features]
      X.iloc[0:5,]
```

```
[67]:   Deg-malig  Age_30-39  Age_40-49  Age_50-59  Age_60-69  Age_70-79  \
0          3          1          0          0          0          0
1          2          0          1          0          0          0
2          2          0          1          0          0          0
3          2          0          0          0          1          0
```

4	2	0	1	0	0	0
---	---	---	---	---	---	---

	Menopause_lt40	Menopause_premeno	Tumor-size_10-14	Tumor-size_15-19	...	\
0	0	1	0	0	...	
1	0	1	0	0	...	
2	0	1	0	0	...	
3	0	0	0	1	...	
4	0	1	0	0	...	

	Inv-nodes_3-5	Inv-nodes_6-8	Inv-nodes_9-11	Node-caps_yes	Breast_right	\
0	0	0	0	0	0	
1	0	0	0	0	1	
2	0	0	0	0	0	
3	0	0	0	0	1	
4	0	0	0	0	1	

	Breast-quad_left_low	Breast-quad_left_up	Breast-quad_right_low	\
0	1	0	0	
1	0	0	0	
2	1	0	0	
3	0	1	0	
4	0	0	1	

	Breast-quad_right_up	IR-Radiat_yes
0	0	0
1	1	0
2	0	0
3	0	0
4	0	0

[5 rows x 31 columns]

```
[68]: #Final Response Set
Y = encoded_BC1_df['Class_recurrence-events']
Y.iloc[0:5,]
```

```
[68]: 0    0
      1    0
      2    0
      3    0
      4    0
      Name: Class_recurrence-events, dtype: uint8
```

```
[69]: #Splitting into training and testing sets
train_X,test_X,train_y,test_y = train_test_split(X,Y,train_size=0.
↪7,random_state=42)
```

[70]: *#Decision Trees using Gini*

```
BC_DT_gini = DecisionTreeClassifier(criterion='gini',max_depth=3)
BC_DT_gini.fit(train_X,train_y)
```

[70]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
max_depth=3, max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort='deprecated',
random_state=None, splitter='best')

[71]: *#Predicting the class of the test set*

```
BC_tree_gini_predict = BC_DT_gini.predict(test_X)
```

[72]: *#AUC, Classification Report and Confusion Matrix*

```
print("AUC: \n", metrics.roc_auc_score(test_y,BC_tree_gini_predict))
print("ROC Curve:\n",metrics.roc_curve(test_y,BC_tree_gini_predict))
print("Classification Report: \n", metrics.
    ↳classification_report(test_y,BC_tree_gini_predict))
print("Confusion Metrics:\n",metrics.
    ↳confusion_matrix(test_y,BC_tree_gini_predict))
```

AUC:

0.5625

ROC Curve:

(array([0. , 0.01785714, 1. , 1.], dtype=float64), array([0. , 0.14285714, 1. , 1.], dtype=float64), array([2, 1, 0]))

Classification Report:


	precision	recall	f1-score	support
0	0.70	0.98	0.81	56
1	0.80	0.14	0.24	28
accuracy			0.70	84
macro avg	0.75	0.56	0.53	84
weighted avg	0.73	0.70	0.62	84

Confusion Metrics:

```
[[55  1]
 [24  4]]
```

[73]: *#Plotting the decision tree based on Gini's index*

```
fig, axes = plt.subplots(nrows = 1,ncols = 1,figsize = (4,4), dpi=300)
fn = encoded_BC1_df[X_features].columns
cn = ["No Recurrence","Recurrence"]
tree.plot_tree(BC_DT_gini,feature_names = fn, class_names=cn,filled = True)
fig.savefig("BC_DT_tree_gini.png")
```



```
[74]: #Decision Trees using Entropy
BC_DT_entropy = DecisionTreeClassifier(criterion='entropy',max_depth=3)
BC_DT_entropy.fit(train_X,train_y)
```

```
[74]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='entropy',
                             max_depth=3, max_features=None, max_leaf_nodes=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, presort='deprecated',
                             random_state=None, splitter='best')
```

```
[75]: #Predicting the class of the test set
BC_tree_entropy_predict = BC_DT_entropy.predict(test_X)
```

```
[76]: #AUC, Classification Report and Confusion Matrix
print("AUC:\n",metrics.roc_auc_score(test_y,BC_tree_entropy_predict))
print("ROC Curve:\n",metrics.roc_curve(test_y,BC_tree_entropy_predict))
print("Classification Report:\n",metrics.
      ↳classification_report(test_y,BC_tree_entropy_predict))
print("Confusion Metrics:\n",metrics.
      ↳confusion_matrix(test_y,BC_tree_entropy_predict))
```

AUC:

0.5357142857142857

ROC Curve:

(array([0. , 0.07142857, 1. , 1.]), array([0. , 0.14285714, 1. , 1.]), array([2, 1, 0]))

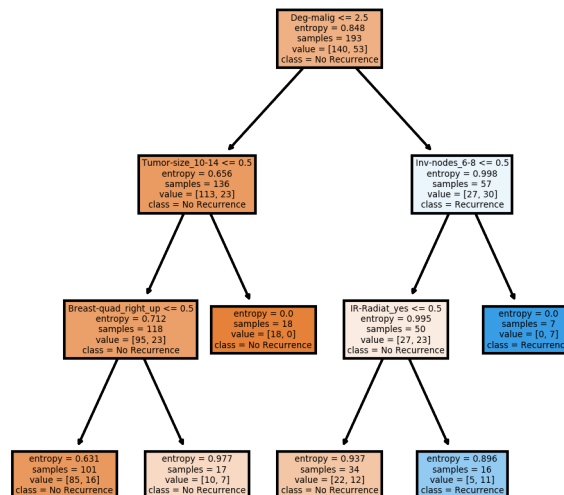
Classification Report:

	precision	recall	f1-score	support
0	0.68	0.93	0.79	56
1	0.50	0.14	0.22	28
accuracy			0.67	84
macro avg	0.59	0.54	0.51	84
weighted avg	0.62	0.67	0.60	84

Confusion Metrics:

```
[[52  4]
 [24  4]]
```

```
[77]: #Ploting the decision tree based on Entropy
fig, axes = plt.subplots(nrows = 1,ncols = 1,figsize = (4,4), dpi=300)
tree.plot_tree(BC_DT_entropy,feature_names = fn, class_names=cn,filled = True);
fig.savefig("BC_DT_tree_entropy.png")
```



```
[78]: #Random Forest Model
from sklearn.ensemble import RandomForestClassifier
BC_tree_RF = RandomForestClassifier(max_depth = 10, n_estimators=10)
```

```
[79]: BC_tree_RF.fit(train_X,train_y)
```

```
[79]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                             criterion='gini', max_depth=10, max_features='auto',
```

```

max_leaf_nodes=None, max_samples=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=10,
n_jobs=None, oob_score=False, random_state=None,
verbose=0, warm_start=False)

```

```

[80]: y_pred_rf = BC_tree_RF.predict(test_X)
      y_pred_rf

```

```

[80]: array([0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0,
            0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0,
            0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
            0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0], dtype=uint8)

```

```

[81]: print("AUC:\n",metrics.roc_auc_score(test_y,y_pred_rf))

```

```

AUC:
0.5714285714285714

```

```

[82]: print("Confusion Matrix\n", (confusion_matrix(test_y,y_pred_rf)))

```

```

Confusion Matrix
[[50  6]
 [21  7]]

```

```

[83]: print("Accuracy:\n", (accuracy_score(test_y,y_pred_rf)*100))

```

```

Accuracy:
67.85714285714286

```

```

[84]: print("Classification Report\n", (classification_report(test_y, y_pred_rf)))

```

```

Classification Report

```

	precision	recall	f1-score	support
0	0.70	0.89	0.79	56
1	0.54	0.25	0.34	28
accuracy			0.68	84
macro avg	0.62	0.57	0.56	84
weighted avg	0.65	0.68	0.64	84

```

[85]: #GRID Search
      from sklearn.model_selection import GridSearchCV
      tuned_parameters = [{'max_depth': [10,15], 'n_estimators': 10,
                             'max_features': ['sqrt',0.2]}]

```

```
BC_tree_clf = RandomForestClassifier()
BC_Grid_clf = GridSearchCV(BC_tree_clf,tuned_parameters,cv=5,scoring='roc_auc')
BC_Grid_clf.fit(train_X,train_y)
```

```
[85]: GridSearchCV(cv=5, error_score=nan,
                estimator=RandomForestClassifier(bootstrap=True, ccp_alpha=0.0,
                                                  class_weight=None,
                                                  criterion='gini', max_depth=None,
                                                  max_features='auto',
                                                  max_leaf_nodes=None,
                                                  max_samples=None,
                                                  min_impurity_decrease=0.0,
                                                  min_impurity_split=None,
                                                  min_samples_leaf=1,
                                                  min_samples_split=2,
                                                  min_weight_fraction_leaf=0.0,
                                                  n_estimators=100, n_jobs=None,
                                                  oob_score=False,
                                                  random_state=None, verbose=0,
                                                  warm_start=False),
                iid='deprecated', n_jobs=None,
                param_grid=[{'max_depth': [10, 15], 'max_features': ['sqrt', 0.2],
                              'n_estimators': [10, 20]}],
                pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                scoring='roc_auc', verbose=0)
```

```
[86]: BC_Grid_clf.best_score_
```

```
[86]: 0.748538961038961
```

```
[87]: BC_Grid_clf.best_params_
```

```
[87]: {'max_depth': 10, 'max_features': 0.2, 'n_estimators': 20}
```

```
[88]: BC_bestTree = ↳RandomForestClassifier(max_depth=10,n_estimators=20,max_features=0.2)
BC_bestTree.fit(train_X,train_y)
```

```
[88]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                             criterion='gini', max_depth=10, max_features=0.2,
                             max_leaf_nodes=None, max_samples=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, n_estimators=20,
                             n_jobs=None, oob_score=False, random_state=None,
                             verbose=0, warm_start=False)
```



```
[89]: y_pred_bestTree = BC_bestTree.predict(test_X)
```

```
[90]: print("Confusion Matrix\n", (confusion_matrix(test_y,y_pred_bestTree)))
```

```
Confusion Matrix
[[52  4]
 [18 10]]
```

```
[91]: print("Accuracy:\n", (accuracy_score(test_y,y_pred_bestTree)*100))
```


```
Accuracy:
73.80952380952381
```

```
[92]: print("Classification Report\n", (classification_report(test_y,
↪y_pred_bestTree)))
```

```
Classification Report
```

	precision	recall	f1-score	support
0	0.74	0.93	0.83	56
1	0.71	0.36	0.48	28
accuracy			0.74	84
macro avg	0.73	0.64	0.65	84
weighted avg	0.73	0.74	0.71	84

```
[93]: #Important Features for the bestTree
import seaborn as sn
feature_rank = pd.DataFrame({'feature': train_X.columns, 'importance':
↪BC_bestTree.feature_importances_})
feature_rank = feature_rank.sort_values('importance', ascending = False)
plt.figure(figsize=(8,6))
sn.barplot(y = 'feature', x = 'importance', data = feature_rank);
```



```
[94]: #Adaboost Classifier
from sklearn.ensemble import AdaBoostClassifier
carseats_adaboost = DecisionTreeClassifier()
BC_tree_Adaboost = AdaBoostClassifier(carseats_adaboost, n_estimators=50)
BC_tree_Adaboost.fit(train_X,train_y)
```

```
[94]: AdaBoostClassifier(algorithm='SAMME.R',
                        base_estimator=DecisionTreeClassifier(ccp_alpha=0.0,
                                                                class_weight=None,
                                                                criterion='gini',
                                                                max_depth=None,
                                                                max_features=None,
                                                                max_leaf_nodes=None,
                                                                min_impurity_decrease=0.0,
                                                                min_impurity_split=None,
                                                                min_samples_leaf=1,
                                                                min_samples_split=2,
                                                                min_weight_fraction_leaf=0.0,
                                                                presort='deprecated',
                                                                random_state=None,
                                                                splitter='best'),
                        learning_rate=1.0, n_estimators=50, random_state=None)
```

```
[95]: y_pred_Adaboost = BC_tree_Adaboost.predict(test_X)
y_pred_Adaboost
```

```
[95]: array([0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0,
          0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0,
          0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0,
          1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0], dtype=uint8)
```

```
[96]: print("AUC:\n",metrics.roc_auc_score(test_y,y_pred_Adaboost))
```

```
AUC:
0.5892857142857143
```

```
[97]: print("Confusion Matrix\n", (confusion_matrix(test_y,y_pred_Adaboost)))
```

```
Confusion Matrix
[[46 10]
 [18 10]]
```

```
[98]: print("Accuracy:\n", (accuracy_score(test_y,y_pred_Adaboost)*100))
```

```
Accuracy:
66.66666666666666
```

```
[99]: print("Classification Report\n", (classification_report(test_y,
→y_pred_Adaboost)))
```

```
Classification Report
```

	precision	recall	f1-score	support
0	0.72	0.82	0.77	56
1	0.50	0.36	0.42	28
accuracy			0.67	84
macro avg	0.61	0.59	0.59	84
weighted avg	0.65	0.67	0.65	84

```
[100]: #Gradient Boosting Classifier
from sklearn.ensemble import GradientBoostingClassifier
BC_tree_gradient = GradientBoostingClassifier(n_estimators = 500, max_depth =
→10)
BC_tree_gradient.fit(train_X,train_y)
```

```
[100]: GradientBoostingClassifier(ccp_alpha=0.0, criterion='friedman_mse', init=None,
learning_rate=0.1, loss='deviance', max_depth=10,
max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
```

```

min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=500,
n_iter_no_change=None, presort='deprecated',
random_state=None, subsample=1.0, tol=0.0001,
validation_fraction=0.1, verbose=0,
warm_start=False)

```

```
[101]: y_pred_gradient = BC_tree_gradient.predict(test_X)
```

```
[102]: print("Confusion Matrix\n", (confusion_matrix(test_y,y_pred_gradient)))
```

```

Confusion Matrix
[[45 11]
 [20  8]]

```

```
[103]: print("Accuracy:\n", (accuracy_score(test_y,y_pred_gradient)*100))
```

```

Accuracy:
63.095238095238095

```

```
[104]: print("Classification Report\n", (classification_report(test_y,
↪y_pred_gradient)))
```

```

Classification Report

```

	precision	recall	f1-score	support
0	0.69	0.80	0.74	56
1	0.42	0.29	0.34	28
accuracy			0.63	84
macro avg	0.56	0.54	0.54	84
weighted avg	0.60	0.63	0.61	84

```
[105]: print("AUC:\n",metrics.roc_auc_score(test_y,y_pred_gradient))
```

```

AUC:
0.5446428571428571

```

```
[ ]:
```